


Training-input effects differ by season even on shared admission targets
Percent change versus the same architecture trained on finalized inputs in all three seasons; three seeds; lower is better



Each arm retrains on 2022-23/2023-24 plus 2025-26 for retrospective evaluation, or plus 2024-25 for forward evaluation.
Errors use only that added season; early two means 2022-23/2023-24. Finalized future labels unchanged. Wednesday archived inputs, B2 lags and final proxies.
C1: pathogen MLP, distance, no covariates. C3: target MLP, neighbors, Kinsa.
C2: target MLP, no spatial; inpatient, wastewater, Kinsa, ILINet, clinical labs and FluSurv.
The two forward columns score the same forecasts: shared admissions receive equal weight; primary admissions have twice ED weight; US20%, states/DC80%.